

MODULE-4

**DATA CLEANING AND SUMMARIZING WITH
DPLYR PACKAGE**

WHAT IS DPLYR?

- The dplyr is a powerful R-package to manipulate, clean and summarize unstructured data. In short, it makes data exploration and data manipulation easy and fast in R.

-
- The package "dplyr" comprises many functions that perform mostly used data manipulation operations such as applying filter, selecting specific columns, sorting data, adding or deleting columns and aggregating data.

How to install and load dplyr package

To install the dplyr package, type the following command.

```
install.packages("dplyr")
```

To load dplyr package, type the command below

```
library(dplyr)
```

Important dplyr Functions

dplyr Function	Description	Equivalent SQL
<code>select()</code>	Selecting columns (variables)	SELECT
<code>filter()</code>	Filter (subset) rows.	WHERE
<code>group_by()</code>	Group the data	GROUP BY
<code>summarise()</code>	Summarise (or aggregate) data	-
<code>arrange()</code>	Sort the data	ORDER BY
<code>join()</code>	Joining data frames (tables)	JOIN
<code>mutate()</code>	Creating New Variables	COLUMN ALIAS

DPLYR PRACTICAL EXAMPLES

EXAMPLE I : SELECTING RANDOM N ROWS

- The **sample_n** function selects random rows from a data frame (or table). The second parameter of the function tells R the number of rows to select.
- Example: `sample_n(mydata,3)`

R

 Copy code

```
# Load the dplyr package
library(dplyr)

# Create a sample dataset
df <- data.frame(
  ID = c(1, 2, 3, 4, 5),
  Category = c("A", "B", "C", "A", "B"),
  Value = c(10, 20, 30, 40, 50)
)

# Use sample_n() to randomly select a specified number of rows
sample_rows <- df %>%
  sample_n(3) # Select 3 random rows from df

print(sample_rows)
```

ID		Category	Value
3	3	C	30
5	5	B	50
1	1	A	10

EXAMPLE 2 : SELECTING RANDOM FRACTION OF ROWS

- The **sample_frac** function returns randomly N% of rows. In the example below, it returns randomly 10% of rows.
- Example: `sample_frac(mydata,0.1)`

R

```
# Use sample_frac() to randomly select a specified fraction of
sample_fraction <- df %>%
  sample_frac(0.5) # Select 50% of the rows from df

print(sample_fraction)
```

ID	Category	Value
5	B	50
4	A	40

EXAMPLE3 :REMOVE DUPLICATE ROWS BASED ON ALL THE VARIABLES

- The **distinct function** is used to eliminate duplicates.
- `x1 = distinct(mydata)`

```
# Use distinct() to get unique rows based on selected columns
unique_rows <- df %>%
  distinct(Category) # Get unique rows based on 'Category' column

print(unique_rows)
```

	Category
1	A
2	B
3	C

SELECT() FUNCTION

- The `select()` function in R is used to choose specific columns from a data frame. It allows you to create a new data frame with only the columns you specify. Here's an example of how to use the `select()` function:

The Following Functions Helps You To Select Variables Based On Their Names.

<code>starts_with()</code>	Starts with a prefix
<code>ends_with()</code>	Ends with a prefix
<code>contains()</code>	Contains a literal string
<code>matches()</code>	Matches a regular expression
<code>num_range()</code>	Numerical range like x01, x02, x03.
<code>one_of()</code>	Variables in character vector.
<code>everything()</code>	All variables.

```
# Create a sample data frame
df <- data.frame(
  ID = c(1, 2, 3),
  Name = c("John", "Jane", "Alice"),
  Age = c(25, 30, 35),
  Gender = c("Male", "Female", "Female"),
  Salary = c(50000, 60000, 70000),
  Country = c("USA", "Canada", "UK")
)

# Select specific columns using select()
selected_df <- select(df, ID, Name, Salary)

print(selected_df)
```

	ID	Name	Salary
1	1	John	50000
2	2	Jane	60000
3	3	Alice	70000

STARTS_WITH(), ENDS_WITH()

- We can also use `select()` with helper functions like `starts_with()`, `ends_with()`, `contains()`, `matches()`, `num_range()`, and `one_of()` to select columns based on specific patterns or conditions.

```
# Select columns based on patterns using select()
selected_pattern <- select(df, starts_with("N"), ends_with("y"))

print(selected_pattern)
```

	Name	Country
1	John	USA
2	Jane	Canada
3	Alice	UK

-
- We can also use the - operator to exclude specific columns from the selection.

```
# Exclude specific columns using select()
excluded_df <- select(df, -Age, -Gender)

print(excluded_df)
```

	ID	Name	Salary	Country
1	1	John	50000	USA
2	2	Jane	60000	Canada
3	3	Alice	70000	UK

CONTAINS()

```
df <- data.frame(  
  ID = c(1, 2, 3),  
  Name = c("John", "Jane", "Alice"),  
  Age = c(25, 30, 35),  
  Gender = c("Male", "Female", "Female"),  
  Salary_2020 = c(50000, 60000, 70000),  
  Salary_2021 = c(55000, 65000, 75000),  
  Country = c("USA", "Canada", "UK")  
)  
  
# Select columns using contains()  
selected_contains <- select(df, contains("Salary"))  
  
print(selected_contains)
```

In this example, `contains("Salary")` is used with `select()` to choose columns that contain the string "Salary". The resulting data frame `selected_contains` includes the columns `Salary_2020`

	Salary_2020	Salary_2021
1	50000	55000
2	60000	65000
3	70000	75000

MATCHES()

	ID	Name	Age	Gender	Salary_2020	Salary_2021	Country
1	1	John	25	Male	50000	55000	USA
2	2	Jane	30	Female	60000	65000	Canada
3	3	Alice	35	Female	70000	75000	UK

```
# Select columns using matches()
selected_matches <- select(df, matches("^S"))

print(selected_matches)
```

	Salary_2020	Salary_2021
1	50000	55000
2	60000	65000
3	70000	75000

Here, `matches("^S")` is used to select columns that start with the letter "S". The resulting data frame `selected_matches` includes the columns `Salary_2020` and `Salary_2021`.

NUM_RANGE()

	ID	Name	Age	Gender	Salary_2020	Salary_2021	Country
1	1	John	25	Male	50000	55000	USA
2	2	Jane	30	Female	60000	65000	Canada
3	3	Alice	35	Female	70000	75000	UK

	Salary_2020	Salary_2021
1	50000	55000
2	60000	65000
3	70000	75000

```
# Select columns using num_range()  
selected_num_range <- select(df, num_range("Salary_", 2020:2021))  
  
print(selected_num_range)
```

In this example, `num_range("Salary_", 2020:2021)` is used to select columns with names starting with "Salary_" followed by the years 2020 and 2021. The resulting data frame `selected_num_range` includes the columns `Salary_2020` and `Salary_2021`.

```
# Create a sample data frame
df <- data.frame(
  Col_1 = c(1, 2, 3),
  Col_2 = c("A", "B", "C"),
  Col_3 = c(4, 5, 6),
  Col_4 = c("D", "E", "F"),
  Col_5 = c(7, 8, 9)
)

# Select columns using num_range()
selected_num_range <- select(df, num_range("Col_", 2:4))

print(selected_num_range)
```

	Col_2	Col_3	Col_4
1	A	4	D
2	B	5	E
3	C	6	F

ONE_OF

```
# Select columns using one_of()
selected_one_of <- select(df, one_of(c("Name", "Country")))

print(selected_one_of)
```

	Name	Country
1	John	USA
2	Jane	Canada
3	Alice	UK

Here, `one_of(c("Name", "Country"))` is used to select columns that match either "Name" or "Country". The resulting data frame `selected_one_of` includes the columns Name and Country

RENAME() FUNCTION

- It is used to change variable name.
 - rename() syntax : `rename(data , new_name = old_name)`
 - data : Data Frame
 - new_name : New variable name you want to keep
 - old_name : Existing Variable Name
- o change variable name.

FILTER() FUNCTION

- In R, the `filter()` function is part of the dplyr package and is used to extract rows from a data frame based on specified conditions. It allows you to subset your data based on logical expressions and filter criteria.
- `filter()` syntax : `filter(data ,)`
- `data` : Data Frame
- `....` : Logical Condition

Example 1: Filter Rows Based On A Single Condition

- # Create a data frame
- `df <- data.frame(ID = 1:5, Name = c("John", "Jane", "Alice", "Bob", "Eve"), Age = c(25, 30, 35, 40, 45))`
- # Filter rows where Age is greater than 30
- `filtered_df <- filter(df, Age > 30)`
- `print(filtered_df)`

	ID	Name	Age
1	1	John	25
2	2	Jane	30
3	3	Alice	35
4	4	Bob	40
5	5	Eve	45

	ID	Name	Age
1	3	Alice	35
2	4	Bob	40
3	5	Eve	45

Example 2: Filter Rows Based On Multiple Conditions

- `filtered_df <- filter(df, Age > 30, startsWith(Name, "J"))`
`print(filtered_df)`

	ID	Name	Age
1	1	John	25
2	2	Jane	30
3	3	Alice	35
4	4	Bob	40
5	5	Eve	45

	ID	Name	Age
1	2	Jane	30

Example 3: Filter Rows Based On Partial String Match

- `filtered_df <- filter(df, grepl("o", Name))`

`print(filtered_df)`

	ID	Name	Age
1	1	John	25
2	2	Jane	30
3	3	Alice	35
4	4	Bob	40
5	5	Eve	45

	ID	Name	Age
1	1	John	25
4	4	Bob	40

Example 4: Filter Rows Based On A Character Condition

- `filtered_df <- filter(df, Name == "Alice")`
- `print(filtered_df)`

	ID	Name	Age
1	1	John	25
2	2	Jane	30
3	3	Alice	35
4	4	Bob	40
5	5	Eve	45

	ID	Name	Age
1	3	Alice	35

Example 5: Filter Rows Based On A Range Of Values

- `filtered_df <- filter(df, Age >= 30 & Age <= 40)`
`print(filtered_df)`

	ID	Name	Age
1	1	John	25
2	2	Jane	30
3	3	Alice	35
4	4	Bob	40
5	5	Eve	45

	ID	Name	Age
1	2	Jane	30
2	3	Alice	35
3	4	Bob	40

FILTER USING IN OPERATOR

- `library(dplyr)`
- `# Create a sample data frame`
- `df <- data.frame(Index = c('A', 'B', 'C', 'D', 'E'), Value = 1:5)`
- `# Filter rows where 'Index' is either 'A' or 'C'`
- `filtered_df <- df %>% filter(Index %in% c('A', 'C'))`
- `# Print the filtered data frame`
- `filtered_df`

	Index	Value
1	A	1
2	B	2
3	C	3
4	D	4
5	E	5

	Index	Value
1	A	1
2	C	3

FILTER USING AND (&) OPERATOR

- # Select rows where 'Index' is 'A' and 'Value' is greater than 2
filtered_df <- df %>% filter(Index == 'A' & Value > 2)
- # Output
- filtered_df

	Index	Value
1	A	1
2	B	2
3	C	3
4	D	4
5	E	5

	Index	Value
1	A	3

FILTER USING OR (|) OPERATOR

- `df <- data.frame(Index = c('A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J'),
Value = 1:10)`
- `# Select rows where 'Index' is 'A' or 'Value' is greater than 5
filtered_df <- df %>% filter(Index == 'A' | Value > 5) # Output
filtered_df`

EXAMPLE

	Index	Value
1	A	1
2	B	2
3	C	3
4	D	4
5	E	5
6	F	6
7	G	7
8	H	8
9	I	9
10	J	10

	Index	Value
1	A	1
2	F	6
3	G	7
4	H	8
5	I	9
6	J	10

-
- `library(dplyr)`
 - `df <- data.frame(Index = c('A', 'B', 'C', 'D', 'E', 'F', 'G', 'H', 'I', 'J'),
Value = 1:10)`
 - `filtered_df <- df %>% filter(!(Index %in% c('C', 'G')))`
`print(filtered_df)`

EXAMPLE

	Index	Value
1	A	1
2	B	2
3	C	3
4	D	4
5	E	5
6	F	6
7	G	7
8	H	8
9	I	9
10	J	10

	Index	Value
1	A	1
2	B	2
4	D	4
5	E	5
6	F	6
8	H	8
9	I	9
10	J	10

SUMMARISE() FUNCTION

- The `summarise()` function in R is a fundamental function from the `dplyr` package used for data manipulation and summarization. It is commonly used in combination with other functions like `group_by()` and `mutate()` to perform data aggregation operations.
- The `summarise()` function allows you to compute summary statistics or create new variables based on grouped data. It takes a dataset or grouped data frame as input and applies a summarizing operation to each group, resulting in a new data frame with the summarized values.

SYNTAX

- `summarise(.data, new_variable = summarizing_operation)`
- `.data`: The input dataset or grouped data frame.
- `new_variable`: The name of the new variable or column to be created.
- `summarizing_operation`: The summarizing operation to be applied to the grouped data.

SOME COMMONLY USED SUMMARIZING OPERATIONS INCLUDE:

- `mean()`: Compute the mean value.
- `median()`: Compute the median value.
- `sum()`: Compute the sum of values.
- `min()`: Find the minimum value.
- `max()`: Find the maximum value.
- `n()`: Count the number of observations.

GROUP_BY()

- The `group_by()` function in R is used to group rows of a dataframe based on one or more variables. Here's an example using the `mtcars` dataset:

- `library(dplyr)`
- `# Grouping the mtcars dataset by the "cyl" variable`
- `grouped_data <- group_by(mtcars, cyl)`
- `# Summarizing the grouped data by calculating the mean of
mpg for each cylinder value summary_data <-
summarize(grouped_data, avg_mpg = mean(mpg))`
- `# Printing the summary data`
- `print(summary_data)`

```
# A tibble: 3 × 2  
  cyl avg_mpg  
  <dbl>   <dbl>  
1     4    26.7  
2     6    19.7  
3     8    15.1
```

MEAN()

```
library(dplyr)

# Create a sample dataset
df <- data.frame(Category = rep(c("A", "B"), each = 5),
                  Value = c(10, 15, 8, 12, 9, 7, 11, 13, 14, 6))

# Compute the mean value for each category
summary_df <- df %>%
  group_by(Category) %>%
  summarise(mean_value = mean(Value))

print(summary_df)
```

	Category	Value
1	A	10
2	A	15
3	A	8
4	A	12
5	A	9
6	B	7
7	B	11
8	B	13
9	B	14
10	B	6

```
# A tibble: 2 × 2
  Category mean_value
  <chr>      <dbl>
1 A         10.8
2 B          8.2
```

MEDIAN(): COMPUTE THE MEDIAN VALUE.

```
library(dplyr)

df <- data.frame(Category = rep(c("A", "B"), each = 5),
                  Value = c(10, 15, 8, 12, 9, 7, 11, 13, 14, 6))

df_summary <- df %>%
  group_by(Category) %>%
  summarize(median_value = median(Value))

print(df_summary)
```

	Category	Value
1	A	10
2	A	15
3	A	8
4	A	12
5	A	9
6	B	7
7	B	11
8	B	13
9	B	14
10	B	6

A tibble: 2 × 2

	Category	median_value
1	A	9
2	B	13

OTHER FUNCTIONS

```
library(dplyr)

# Example dataframe
df <- data.frame(Category = rep(c("A", "B"), each = 5),
                  Value = c(10, 15, 8, 12, 9, 7, 11, 13, 14, 6))

# Calculate sum, minimum, maximum, and count for each category
df_summary <- df %>%
  group_by(Category) %>%
  summarize(sum_value = sum(Value),
            min_value = min(Value),
            max_value = max(Value),
            count = n())

print(df_summary)
```

	Category	Value
1	A	10
2	A	15
3	A	8
4	A	12
5	A	9
6	B	7
7	B	11
8	B	13
9	B	14
10	B	6

	Category	sum_value	min_value	max_value	count
	<chr>	<dbl>	<dbl>	<dbl>	<int>
1	A	54	8	15	5
2	B	51	6	14	5

ARRANGE() FUNCTION

It Is Used To Sort The Data

- In R, the `arrange()` function is used to reorder rows in a dataframe based on one or more columns. It allows you to sort the data in ascending or descending order. The `arrange()` function is part of the `dplyr` package and is commonly used for data manipulation and exploratory data analysis.

SYNTAX

- `arrange(data, column1, column2, ...)`
- `data`: The dataframe or tibble you want to arrange.
- `column1, column2, ...`: The columns to sort the data by. You can specify multiple columns to sort by. By default, the sorting is done in ascending order.

EXAMPLE

```
library(dplyr)

# Example dataframe
df <- data.frame(Name = c("Alice", "Bob", "John", "Jane"),
                  Age = c(25, 30, 35, 40),
                  Score = c(80, 95, 70, 85))

# Sort the dataframe by Age in ascending order
df_sorted <- arrange(df, Age)

print(df_sorted)
```

	Name	Age	Score
1	Alice	25	80
2	Bob	30	95
3	John	35	70
4	Jane	40	85

	Name	Age	Score
1	Alice	25	80
2	Bob	30	95
3	John	35	70
4	Jane	40	85

DO()

- In R, the `do()` function is part of the `dplyr` package and allows you to perform custom operations or transformations on grouped data. It is often used in combination with other `dplyr` functions like `group_by()` and `summarize()`. Here's an example to illustrate its usage:

```
library(dplyr)

# Creating a dataframe
df <- data.frame(Category = c("A", "A", "B", "B", "C", "C"),
                  Value = c(10, 15, 8, 12, 9, 7))

# Grouping the dataframe by Category and applying custom function using do()
result <- df %>%
  group_by(Category) %>%
  do(custom_summary = sum(.$Value) / mean(.$Value))

# Printing the result
print(result)
```

```
# A tibble: 3 × 2
# Groups:   Category [3]
  Category custom_summary
  <chr>         <dbl>
1 A             2.33
2 B             2.33
3 C             2.33
```

In this example, we grouped the dataframe `df` by the "Category" column using `group_by()`. Then, we used `do()` to apply a custom function that calculates the sum of "Value" divided by the mean of "Value" within each group. The resulting dataframe `result` contains the grouped categories and the corresponding custom summary value for each group.

MUTATE() FUNCTION

- In R, the `mutate()` function is part of the `dplyr` package and is used to create or modify variables (columns) within a dataframe. It allows you to perform calculations or transformations on existing variables and create new variables based on those calculations. Here's an example to illustrate its usage:


```
library(dplyr)

# Creating a dataframe
df <- data.frame(Name = c("Alice", "Bob", "John", "Jane"),
                  Age = c(25, 30, 35, 40),
                  Score = c(80, 95, 70, 85))

# Adding a new variable using mutate()
df <- df %>%
  mutate(Age_Group = ifelse(Age < 30, "Young", "Old"))

# Printing the modified dataframe
print(df)
```

	Name	Age	Score	Age_Group
1	Alice	25	80	Young
2	Bob	30	95	Old
3	John	35	70	Old
4	Jane	40	85	Old

JOIN() FUNCTION

In R, the `join()` function is used to combine two or more data frames based on a common variable or set of variables. It is part of the `dplyr` package and provides several types of joins, including inner join, left join, right join, and full join. Here's an overview of the different join functions available in `dplyr`:

1. Inner Join (`inner_join()`): Returns only the rows that have matching values in both data frames.
2. Left Join (`left_join()`): Returns all the rows from the left data frame and the matched rows from the right data frame. If there is no match, it includes NA values for the right data frame columns.
3. Right Join (`right_join()`): Returns all the rows from the right data frame and the matched rows from the left data frame. If there is no match, it includes NA values for the left data frame columns.
4. Full Join (`full_join()`): Returns all the rows from both data frames. If there is no match, it includes NA values for the unmatched columns.

```
df1 <- data.frame(ID = c(1, 2, 3),
                  Name = c("Alice", "Bob", "John"))

df2 <- data.frame(ID = c(2, 3, 4),
                  Age = c(25, 30, 35))

# Inner join
inner_result <- inner_join(df1, df2, by = "ID")

# Left join
left_result <- left_join(df1, df2, by = "ID")

# Right join
right_result <- right_join(df1, df2, by = "ID")

# Full join
full_result <- full_join(df1, df2, by = "ID")

# Printing the join results
print(inner_result)
print(left_result)
print(right_result)
print(full_result)
```

	ID	Name	Age
--	----	------	-----

1	2	Bob	25
---	---	-----	----

2	3	John	30
---	---	------	----

	ID	Name	Age
--	----	------	-----

1	1	Alice	NA
---	---	-------	----

2	2	Bob	25
---	---	-----	----

3	3	John	30
---	---	------	----

	ID	Name	Age
--	----	------	-----

1	2	Bob	25
---	---	-----	----

2	3	John	30
---	---	------	----

3	4	NA	35
---	---	----	----

	ID	Name	Age
--	----	------	-----

1	1	Alice	NA
---	---	-------	----

2	2	Bob	25
---	---	-----	----

3	3	John	30
---	---	------	----

4	4	NA	35
---	---	----	----

INTERSECT() AND UNION()

- In R, the `intersect()` and `union()` functions are used to perform set operations on vectors or data frames. Here's an explanation of each function:
 1. `intersect(x, y)`: Returns a vector or data frame containing the common elements between two input vectors or data frames `x` and `y`. The resulting elements are unique and appear in the order of `x`. If the input arguments are data frames, the columns must have the same names and types.
 2. `union(x, y)`: Returns a vector or data frame containing all the unique elements from both input vectors or data frames `x` and `y`. The resulting elements appear in the order they first appear in the combined vectors. If the input arguments are data frames, the columns must have the same names and types.

EXAMPLE

```
# Example using intersect()
vector1 <- c(1, 2, 3, 4, 5)
vector2 <- c(4, 5, 6, 7, 8)

intersect_result <- intersect(vector1, vector2)
print(intersect_result) # Output: 4 5

# Example using union()
vector3 <- c(3, 4, 5, 6, 7)

union_result <- union(vector1, vector3)
print(union_result) # Output: 1 2 3 4 5 6 7
```


SETDIFF()

- In R, the `setdiff()` function is used to find the set difference between two vectors or data frames. It returns the elements from the first input that are not present in the second input.
- Here's the syntax of the `setdiff()` function:

EXAMPLE

```
# Example using setdiff()
vector1 <- c(1, 2, 3, 4, 5)
vector2 <- c(4, 5, 6, 7, 8)

setdiff_result <- setdiff(vector1, vector2)
print(setdiff_result) # Output: 1 2 3
```

SLICE()



In R, the `slice()` function is used to extract specific rows from a data frame or tibble based on their positions. It allows you to select rows by specifying the row indices or a range of row indices.

EXAMPLE

- `df <- data.frame(ID = 1:5, Name = c("John", "Jane", "Alice", "Bob", "Eve"), Age = c(25, 30, 35, 40, 45))`
- `# Extract the second and fourth rows`
- `result <- slice(df, c(2, 4))`
- `# Print the result`
- `print(result)`

	ID	Name	Age
1	2	Jane	30
2	4	Bob	40

INTERSECT() AND UNION()

- In R, the `intersect()` and `union()` functions are used to perform set operations on vectors or data frames. Here's an explanation of each function:
 1. `intersect(x, y)`: Returns a vector or data frame containing the common elements between two input vectors or data frames `x` and `y`. The resulting elements are unique and appear in the order of `x`. If the input arguments are data frames, the columns must have the same names and types.
 2. `union(x, y)`: Returns a vector or data frame containing all the unique elements from both input vectors or data frames `x` and `y`. The resulting elements appear in the order they first appear in the combined vectors. If the input arguments are data frames, the columns must have the same names and types.